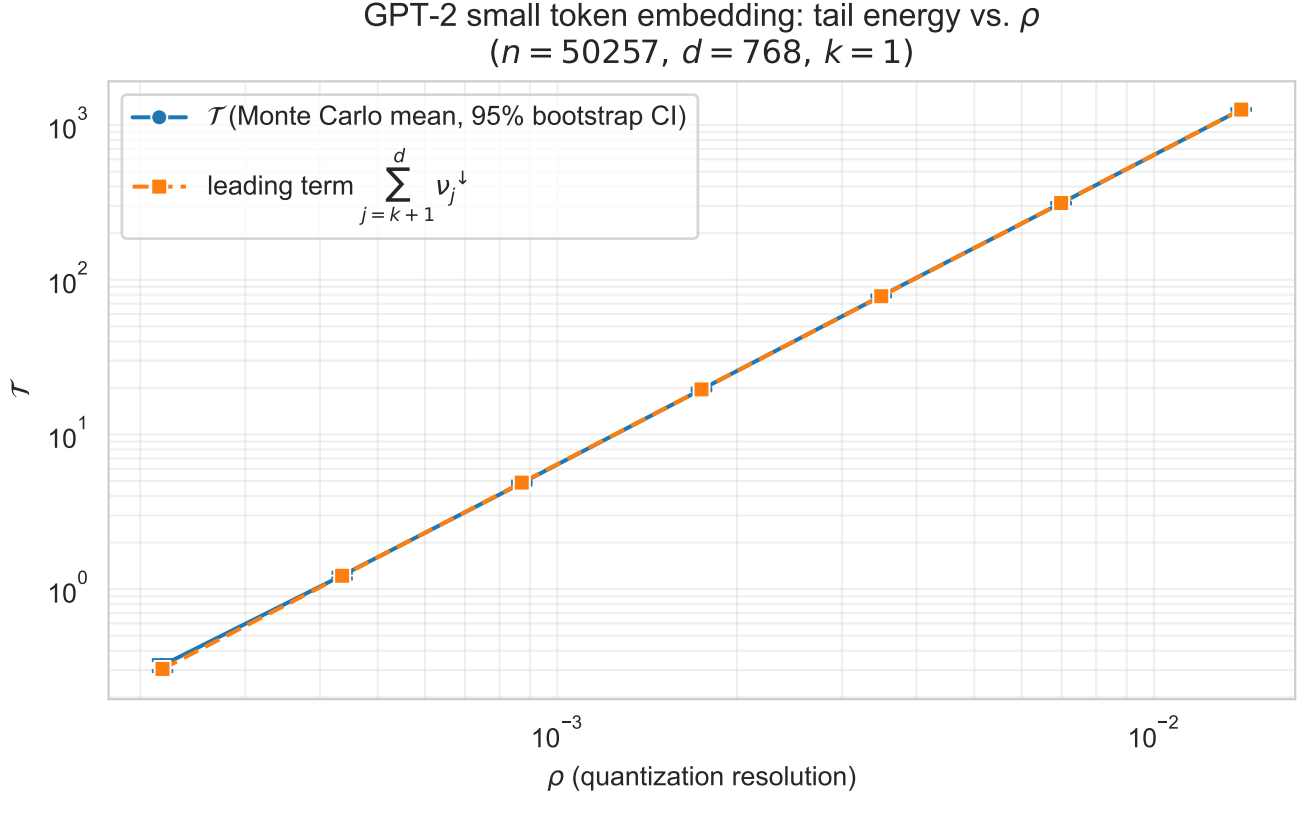
GPT-2 small token embedding: tail energy vs. ρ
($n = 50257$, $d = 768$, $k = 1$)
$\mathcal{T}$ (Monte Carlo mean, 95% bootstrap CI)
leading term $\sum_{j=k+1}^{d} \nu_j^{\downarrow}$
10^{3}
10^{2}
10^{1}
10^{0}
$\mathcal{T}$
10^{-3}
10^{-2}
ρ (quantization resolution)